

```

In [1]: # to handle datasets
import pandas as pd
import numpy as np

# for plotting
import matplotlib.pyplot as plt
%matplotlib inline

# to divide train and test set
from sklearn.model_selection import train_test_split

# feature scaling
from sklearn.preprocessing import StandardScaler

# for tree binarisation
from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import cross_val_score

# to build the models
from sklearn.linear_model import LinearRegression, Lasso
from sklearn.ensemble import RandomForestRegressor
from sklearn.svm import SVR
import xgboost as xgb

# to evaluate the models
from sklearn.metrics import mean_squared_error

pd.pandas.set_option('display.max_columns', None)

import warnings
warnings.filterwarnings('ignore')

```

```

In [2]: # load dataset
df = pd.read_csv('dat102_cleaned.csv')
print(df.shape)
df.head()

```

```
(1280, 47)
```

Out[2]:

	row_id	country_code	year	agricultural_land_area	percentage_of_arable_land_eq
0	1248	1881055	2000	12364.30652	27.873476
1	887	1881055	2001	12292.06398	27.627966
2	969	1881055	2002	12402.83412	27.724205
3	993	1881055	2003	12244.51139	28.454720
4	591	1881055	2004	12171.79164	28.514809

```
In [3]: # Load the dataset for submission (the one on which our model will be evaluated by Kaggle)
# it contains exactly the same variables, but not the target

submission = pd.read_csv('dat102testcleaned.csv')
submission.head()
```

Out[3]:

	row_id	country_code	year	agricultural_land_area	percentage_of_arable_land_eq
0	0	c3ce4bf	2012	15020.598000	99.417846
1	1	1af00b8	2009	11899.695840	54.339839
2	2	dbc74a3	2015	573171.401200	33.070081
3	3	c72199a	2003	157313.946300	4.793346
4	4	8191231	2001	340.049453	33.070081

```
In [4]: # find categorical variables
categorical = [var for var in df.columns if df[var].dtype=='O']
print('There are {} categorical variables'.format(len(categorical)))

There are 1 categorical variables
```

```
In [5]: # find numerical variables
numerical = [var for var in df.columns if df[var].dtype!='O']
print('There are {} numerical variables'.format(len(numerical)))

There are 46 numerical variables
```

```
In [6]: #explore any missing years for train set
pd.set_option('display.max_rows',100)
table1 = pd.crosstab(df['country_code'],df['year'])
display(table1)
```

[illegible]

```
In [7]: #explore any missing years for test set
pd.set_option('display.max_rows',100)
table2 = pd.crosstab(submission['country_code'],submission['year'])
display(table2)
```

[illegible]

```
In [8]: # let's visualise the values of the discrete variables
discrete = []
for var in numerical:
    if len(df[var].unique())<20:
        print(var, ' values: ', df[var].unique())
        discrete.append(var)

print('There are {} discrete variables'.format(len(discrete)))

year  values:  [2000 2001 2002 2003 2004 2005 2006 2007 2008 2009
2010 2011 2012 2013
2014 2015]
There are 1 discrete variables
```

```
In [9]: continuous = [var for var in numerical if var not in discrete and var not in ['row_id', 'prevalence_of_overnourishment']]
continuous
```

```
Out[9]: ['agricultural_land_area',
'percentage_of_arable_land_equipped_for_irrigation',
'cereal_yield',
'droughts_floods_extreme_temps',
'forest_area',
'total_land_area',
'fertility_rate',
'life_expectancy',
'rural_population',
'total_population',
'urban_population',
'population_growth',
'avg_value_of_food_production',
'cereal_import_dependency_ratio',
'food_imports_as_share_of_merch_exports',
'gross_domestic_product_per_capita_ppp',
'imports_of_goods_and_services',
'inequality_index',
'net_oda_received_percent_gni',
'net_oda_received_per_capita',
'tax_revenue_share_gdp',
'trade_in_services',
'per_capita_food_production_variability',
'per_capita_food_supply_variability',
'adult_literacy_rate',
'school_enrollment_rate_female',
'school_enrollment_rate_total',
'avg_supply_of_protein_of_animal_origin',
'caloric_energy_from_cereals_roots_tubers',
'access_to_improved_sanitation',
'access_to_improved_water_sources',
'anemia_prevalence',
'obesity_prevalence',
'open_defecation',
'hiv_incidence',
'rail_lines_density',
'access_to_electricity',
'co2_emissions',
'unemployment_rate',
'total_labor_force',
'military_expenditure_share_gdp',
'proportion_of_seats_held_by_women_in_gov',
'political_stability']
```




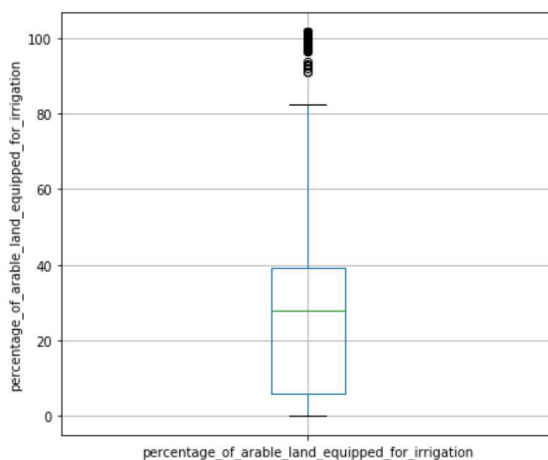








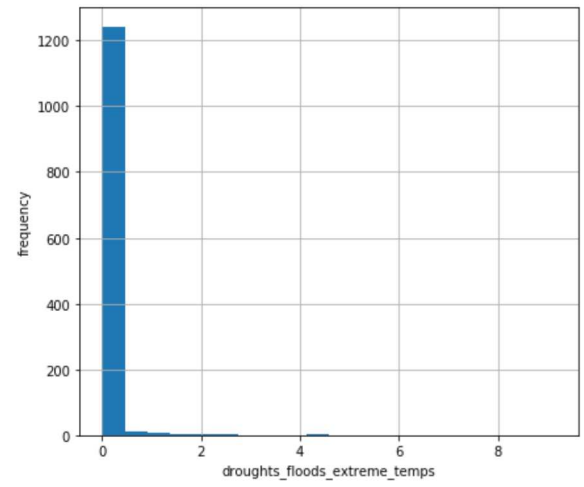
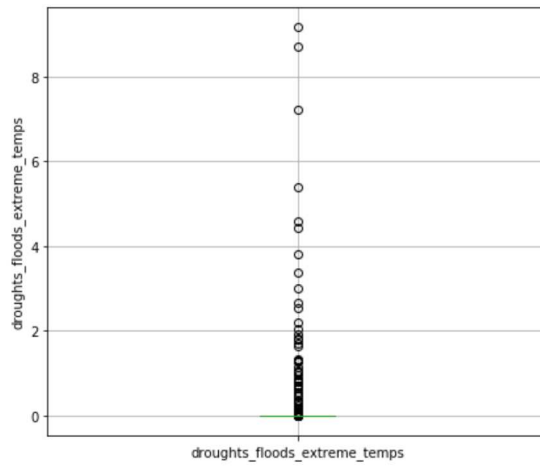
```
In [10]: # let's make boxplots to visualise outliers in the continuous variables
# and histograms to get an idea of the distribution







for var in continuous:
    plt.figure(figsize=(15,6))
    plt.subplot(1, 2, 1)
    fig = df.boxplot(column=var)
    fig.set_title('')
    fig.set_ylabel(var)







    plt.subplot(1, 2, 2)
    fig = df[var].hist(bins=20)
    fig.set_ylabel('frequency')
    fig.set_xlabel(var)


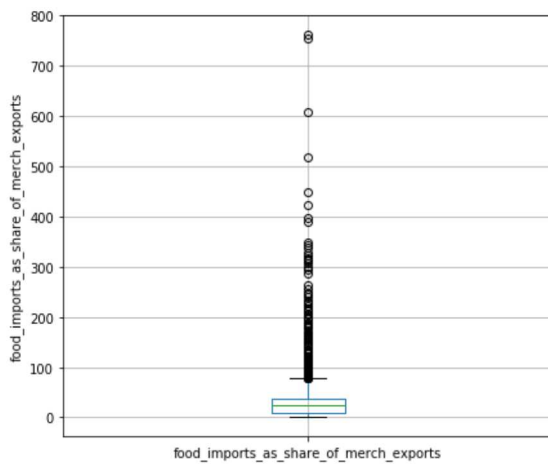




plt.show()
```


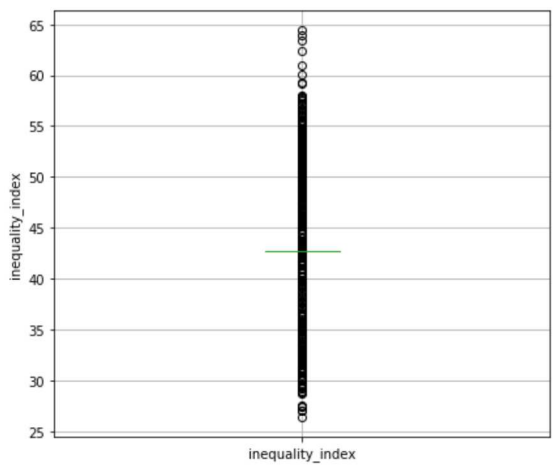




















































































```
In [11]: # outliers in discrete variables
for var in discrete:
    print(df[var].value_counts() / np.float(len(df)))
    print()
```

```
2015    0.0625
2014    0.0625
2013    0.0625
2012    0.0625
2011    0.0625
2010    0.0625
2009    0.0625
2008    0.0625
2007    0.0625
2006    0.0625
2005    0.0625
2004    0.0625
2003    0.0625
2002    0.0625
2001    0.0625
2000    0.0625
Name: year, dtype: float64
```

```
In [12]: for var in categorical:
          print(var, ' contains ', len(df[var].unique()), ' labels')

country_code  contains  80  labels
```

Separate train and test set


```
In [13]: # Let's separate into train and test set

X_train, X_test, y_train, y_test = train_test_split(df, df.prevalence_of_undernourishment, test_size=0.2, random_state=0)
X_train.shape, X_test.shape

Out[13]: ((1024, 47), (256, 47))
```

Outliers in Numerical variables

In order to tackle outliers and skewed distributions at the same time, I suggested I would do discretisation. And in order to find the optimal buckets automatically, I would use decision trees to find the buckets for me.

```
In [14]: def tree_binariser(var):
    score_ls = [] # here I will store the mse

    for tree_depth in [1,2,3,4]:
        # call the model
        tree_model = DecisionTreeRegressor(max_depth=tree_depth)

        # train the model using 3 fold cross validation
        scores = cross_val_score(tree_model, X_train[var].to_frame(), y_train, cv=3, scoring='neg_mean_squared_error')
        score_ls.append(np.mean(scores))

    # find depth with smallest mse
    depth = [1,2,3,4][np.argmax(score_ls)]
    #print(score_ls, np.argmax(score_ls), depth)

    # transform the variable using the tree
    tree_model = DecisionTreeRegressor(max_depth=depth)
    tree_model.fit(X_train[var].to_frame(), X_train.prevalence_of_undernourishment)
    X_train[var] = tree_model.predict(X_train[var].to_frame())
    X_test[var] = tree_model.predict(X_test[var].to_frame())
    submission[var] = tree_model.predict(submission[var].to_frame())

In [15]: for var in continuous:
    tree_binariser(var)
```

```
In [16]: X_train[continuous].head()
```

```
Out[16]:
```

	agricultural_land_area	percentage_of_arable_land_equipped_for_irrigation	cereals
308	15.467345	14.406020	12.80
295	15.467345	14.406020	7.866
915	15.467345	14.406020	12.80
465	26.980645	18.339051	17.40
298	15.467345	14.406020	7.866

```
In [17]: for var in continuous:
          print(var, len(X_train[var].unique()))
```

```
agricultural_land_area 16
percentage_of_arable_land_equipped_for_irrigation 4
cereal_yield 8
droughts_floods_extreme_temps 2
forest_area 16
total_land_area 16
fertility_rate 4
life_expectancy 8
rural_population 16
total_population 16
urban_population 8
population_growth 4
avg_value_of_food_production 8
cereal_import_dependency_ratio 16
food_imports_as_share_of_merch_exports 4
gross_domestic_product_per_capita_ppp 8
imports_of_goods_and_services 8
inequality_index 2
net_oda_received_percent_gni 4
net_oda_received_per_capita 8
tax_revenue_share_gdp 2
trade_in_services 8
per_capita_food_production_variability 4
per_capita_food_supply_variability 2
adult_literacy_rate 2
school_enrollment_rate_female 4
school_enrollment_rate_total 4
avg_supply_of_protein_of_animal_origin 8
caloric_energy_from_cereals_roots_tubers 8
access_to_improved_sanitation 4
access_to_improved_water_sources 8
anemia_prevalence 4
obesity_prevalence 8
open_defecation 4
hiv_incidence 8
rail_lines_density 8
access_to_electricity 8
co2_emissions 16
unemployment_rate 2
total_labor_force 8
military_expenditure_share_gdp 2
proportion_of_seats_held_by_women_in_gov 15
political_stability 4
```

Feature scaling

```
In [18]: X_train.describe()
```

Out[18]:

	row_id	year	agricultural_land_area	percentage_of_arable_land_
count	1024.000000	1024.000000	1024.000000	1024.000000
mean	699.700195	2007.336914	16.422620	16.422620
std	404.644165	4.598309	7.244404	3.614167
min	0.000000	2000.000000	4.617564	14.406020
25%	349.750000	2003.000000	10.044163	14.406020
50%	690.500000	2007.000000	15.467345	14.406020
75%	1039.250000	2011.000000	20.664900	18.339051
max	1400.000000	2015.000000	40.541780	31.866695

```
In [19]: training_vars = [var for var in X_train.columns if var not in ['row_id', 'country_code', 'year', 'prevalence_of_undernourishment']]
```

```
In [20]: # fit scaler
scaler = StandardScaler() # create an instance
scaler.fit(X_train[training_vars]) # fit the scaler to the train set for later use
```

Out[20]: StandardScaler(copy=True, with_mean=True, with_std=True)

Machine Learning algorithm building

xgboost

```
In [21]: xgb_model = xgb.XGBRegressor(max_depth=32, learning_rate=0.01, n_estimators=1000, silent=True, )

eval_set = [(X_test[training_vars], y_test)]
xgb_model.fit(X_train[training_vars], y_train, eval_set=eval_set, verbose=False)

pred = xgb_model.predict(X_train[training_vars])
print('xgb train mse: {}'.format(mean_squared_error(y_train, pred)))
pred = xgb_model.predict(X_test[training_vars])
print('xgb test mse: {}'.format(mean_squared_error(y_test, pred)))

xgb train mse: 0.042008455030963364
xgb test mse: 11.23932270180297
```

Random Forests

```
In [22]: rf_model = RandomForestRegressor(n_estimators=2000,max_depth=32,crit
erion="mse",random_state=1234)
rf_model.fit(X_train[training_vars], y_train)

pred = rf_model.predict(X_train[training_vars])
print('rf train mse: {}'.format(mean_squared_error(y_train, pred)))
pred = rf_model.predict(X_test[training_vars])
print('rf test mse: {}'.format(mean_squared_error(y_test, pred)))

rf train mse: 0.8891796707427881
rf test mse: 14.013620044019174
```

Support vector machine

```
In [23]: SVR_model = SVR()
SVR_model.fit(scaler.transform(X_train[training_vars]), y_train)

pred = SVR_model.predict(scaler.transform(X_train[training_vars]))
print('SVR train mse: {}'.format(mean_squared_error(y_train, pred)))
pred = SVR_model.predict(scaler.transform(X_test[training_vars]))
print('SVR test mse: {}'.format(mean_squared_error(y_test, pred)))

SVR train mse: 37.01011864816445
SVR test mse: 59.79989098398971
```

Regularised linear regression

```
In [24]: lin_model = Lasso(random_state=2909)
lin_model.fit(scaler.transform(X_train[training_vars]), y_train)

pred = lin_model.predict(scaler.transform(X_train[training_vars]))
print('linear train mse: {}'.format(mean_squared_error(y_train, pred
)))
pred = lin_model.predict(scaler.transform(X_test[training_vars]))
print('linear test mse: {}'.format(mean_squared_error(y_test, pred))
)

linear train mse: 30.79952937572478
linear test mse: 53.28622729247063
```

Submission to DataDriven

```
In [ ]: pred_ls = [] #combine xgboost and random forest results
        for model in [xgb_model, rf_model]:
            pred_ls.append(pd.Series(model.predict(submission[training_vars]
            )))

        #pred = SVR_model.predict(scaler.transform(submission[training_vars]
        ))
        #pred_ls.append(pd.Series(pred))

        #pred = lin_model.predict(scaler.transform(submission[training_vars]
        ))
        #pred_ls.append(pd.Series(pred))

        final_pred = pd.concat(pred_ls, axis=1).mean(axis=1)
```

```
In [ ]: temp = pd.concat([submission.row_id, final_pred], axis=1)
        temp.columns = ['row_id', 'prevalence_of_overnourishment']
        temp.head()
```

```
In [ ]: temp.to_csv('resultxgrf.csv', index=False)
```

```
In [ ]: #For XGBoost
        pred_ls5 = []
        for model in [xgb_model]:
            pred_ls5.append(pd.Series(model.predict(submission[training_vars]
            )))
        final_pred = pd.concat(pred_ls5, axis=1).mean(axis=1)
```

```
In [ ]: temp5 = pd.concat([submission.row_id, final_pred], axis=1)
        temp5.columns = ['row_id', 'prevalence_of_overnourishment']
        temp5.head()
```

```
In [ ]: #pd.pandas.set_option('display.max_rows', None)
        #display(temp)
```

```
In [ ]: temp5.to_csv('resultxg.csv', index=False)
```

```
In [ ]: #For Random Forest
        pred_ls2 = []
        for model in [rf_model]:
            pred_ls2.append(pd.Series(model.predict(submission[training_vars]
            )))
        final_pred = pd.concat(pred_ls2, axis=1).mean(axis=1)
```

```
In [ ]: temp2 = pd.concat([submission.row_id, final_pred], axis=1)
        temp2.columns = ['row_id', 'prevalence_of_overnourishment']
        temp2.head()
```

```
In [ ]: temp2.to_csv('resulttrf.csv', index=False)
```

```
In [ ]: #For Support Vector Machine
pred_ls3 = []
for model in [SVR_model]:
    pred_ls3.append(pd.Series(model.predict(submission[training_vars
])))
final_pred = pd.concat(pred_ls3, axis=1).mean(axis=1)
```

```
In [ ]: temp3 = pd.concat([submission.row_id, final_pred], axis=1)
temp3.columns = ['row_id', 'prevalence_of_undernourishment']
temp3.head()
```

```
In [ ]: temp3.to_csv('resultsvm.csv', index=False)
```

```
In [ ]: #For Linear regression
pred_ls4 = []
for model in [lin_model]:
    pred_ls4.append(pd.Series(model.predict(submission[training_vars
])))
final_pred = pd.concat(pred_ls4, axis=1).mean(axis=1)
```


```
In [ ]: temp4 = pd.concat([submission.row_id, final_pred], axis=1)
temp4.columns = ['row_id', 'prevalence_of_undernourishment']
temp4.head()
```

```
In [ ]: temp4.to_csv('resultslin.csv', index=False)
```

Feature importance


```
In [25]: importance = pd.Series(rf_model.feature_importances_)
importance.index = training_vars
importance.sort_values(inplace=True, ascending=False)
importance.plot.bar(figsize=(18,6))
```

Out[25]: <matplotlib.axes._subplots.AxesSubplot at 0xd6619ea20>




```
In [26]: importance = pd.Series(xgb_model.feature_importances_)
importance.index = training_vars
importance.sort_values(inplace=True, ascending=False)
importance.plot.bar(figsize=(18,6))
```

Out[26]: <matplotlib.axes._subplots.AxesSubplot at 0xd66a66cc0>




```
In [27]: importance = pd.Series(np.abs(lin_model.coef_.ravel()))
importance.index = training_vars
importance.sort_values(inplace=True, ascending=False)
importance.plot.bar(figsize=(18,6))
```

Out[27]: <matplotlib.axes._subplots.AxesSubplot at 0xd66b1dc88>



```
In [28]: importance = pd.Series(np.abs(SVR_model.coef_.ravel()))
importance.index = training_vars
importance.sort_values(inplace=True, ascending=False)
importance.plot.bar(figsize=(18,6))
```

```
-----
AttributeError                                Traceback (most recent c
all last)
```

```
<ipython-input-28-26ab1fde4b19> in <module>
```

```
----> 1 importance = pd.Series(np.abs(SVR_model.coef_.ravel()))
      2 importance.index = training_vars
      3 importance.sort_values(inplace=True, ascending=False)
      4 importance.plot.bar(figsize=(18,6))
```

```
~\Anaconda3\lib\site-packages\sklearn\svm\base.py in coef_(self)
```

```
    482     def coef_(self):
    483         if self.kernel != 'linear':
--> 484             raise AttributeError('coef_ is only available
when using a '
    485                                 'linear kernel')
    486
```

```
AttributeError: coef_ is only available when using a linear kernel
```